






The edit map, one average vector per position

One panel per model: these are five different maps and are never pooled. Each arrow is the mean of every edit observed leaving that position. Dot size is how many edits the mean is over. Grey = mean smaller than the rater noise on a mean of that many edits. Arrows drawn 1.0x.



How to read a panel

Each dot is a position some constitution occupied on these three axes. The arrow from it is the average of every edit that model made starting from there.

A long arrow means the model reliably moves documents at that position in one direction. A grey arrow means the average is smaller than the judge disagreeing with itself, so nothing can be concluded from it.

Big dots are better estimates: the mean is over more edits. Positions are integer ratings, so distinct constitutions that score the same on these three axes share a dot and are averaged together - they may differ on the other nine axes.

Green marks the seed each chain started from.

— mean above rater noise — mean within rater noise • position (1 edit) • position (many edits) • seed (model spec)